


Linux Redirection and Pipes

Outline

- STDIN (Standard Input)
- STDOUT (Standard Output)
- STDERR (Standard Error)
- Redirection
- Piping Commands




STDIN, STDOUT and STDERR

- File Descriptors

- When the program is started, the shell gives all programs access to three “pointers” called

- standard input (**STDIN**),
- standard output (**STDOUT**),
- and standard error (**STDERR**).



- Normally **STDIN** is connected to the **keyboard**, while **STDOUT** and **STDERR** are connected to the **screen**.
- However, with **redirection** (**>** , **>>** , **<**) you can change where a program gets its' input from, and where it sends its output to.

STDIN, STDOUT and STDERR

- The redirection operator “>” lets you specify a file name to receive the STDOUT (output) of the command, while the “<” operator lets you connect STDIN to a file containing the input.

Name	Abbreviation	File Descriptor	Standard Device
Standard Input	<i>stdin</i>	0	Keyboard
Standard Output	<i>stdout</i>	1	Console
Standard Error	<i>stderr</i>	2	Console

Examples of Redirection

<code>cat < file1.txt</code>	<i>Redirects from a file</i> < tells cat to take the input from the file instead of the keyboard
<code>ls</code>	standard output for ls is the terminal
<code>ls > lsout.txt</code>	Redirect Output of ls to a file
<code>ls >> lsout.txt</code>	Redirect and append output to a file

Input and Output Redirection

- Default Standard Output is sent to the terminal.
- Redirect Standard Output to a file:
 - `ls -al > martin.txt`
 - `ls -al 1> martin.txt` [This is identical to the command above]
 - Please note that `>` and `1>` will overwrite anything in the `martin.txt` file if it exists, or create the file if it does not exist.
 - `cat martin.txt`
- Redirect standard input to a command.
 - `cat < martin.txt`
 - `cat 0< martin.txt`
- Count the number of lines in a file
 - `wc -l martin.txt`
 - `wc -l < martin.txt`

Error Redirection

- Default Standard Error is sent to the terminal
 - `rm martin.txt`
 - `cat martin.txt`
 - If `martin.txt` does not exist we get an error message like...
 - `cat: martin.txt: No such file or directory`
- Redirecting Standard Error to a file:
 - `cat martin.txt 2> errorfile.txt`
 - `cat errorfile.txt`
 - `cat: martin.txt: No such file or directory`
- Redirecting both Standard Output and Standard Error to the same file:
 - `cat ~/.profile martin.txt > logfile.txt 2>&1`
 - [the `.profile` file exists; the `martin.txt` file does not exist]
 - `cat ~/.profile martin.txt &> logfile.txt`

Redirect Output or Error Stream to Append

- Redirect Standard Output to append (using >>)
 - `ls -al > martin.txt`
 - `echo "Thank you" >> martin.txt`
- Redirect Standard Error to append (using >>)
 - `rm martin.txt`
 - `cat martin.txt`
 - `cat martin.txt 2> error.txt`
 - `echo "I don't like errors" 2>> errors.txt`

Practice

- `echo "b" > file.txt`
 - `echo "a" >> file.txt`
 - `echo "d" >> file.txt`
 - `echo "c" >> file.txt`
 - `cat file.txt`
 - `sort < file.txt`
 - **`sort < file.txt > sortedfile.txt`**
 - `cat sortedfile.txt`
-
- **In the bold command above, we have redirected both the standard input and standard output streams to a file.**

Redirection to standard Output or standard error

- Redirect to Standard Output
 - `ls -al >&1`
 - `echo "Thank you" >&1` [this is the default behaviour anyway]
- Redirect to Standard Error
 - `rm martin.txt`
 - `cat martin.txt >&2`
 - `echo "I don't like errors" >&2`
- Redirecting BOTH to standard output AND to standard error can be useful when piping commands
[piping commands are coming up shortly]

Very useful redirection

- Discard the standard output of a command
 - Syntax: **command > /dev/null**
 - The special file `/dev/null` discards all data written to it
 - `ls -al > /dev/null`
- Redirect both to standard output AND to a file.
 - Syntax: **command | tee file**
 - `ls -al | tee martin.txt`

Piping Commands

- Send STDOUT of one process to STDIN of another process
 - Syntax: **command-1 | command-2 | ... | command-N**
 - `ls -al | grep file`
 - `ls -al /etc | grep x1`
 - `ls -al /etc | grep -i x1`
- **grep**: print lines that match patterns
 - The **-i** option means ignore case sensitivity.
 - Very useful for searching contents of files and standard output.

Useful Piping Commands

- **grep:** print lines that match patterns
 - The **-i** option means match both upper and lower case
 - The **-c** option means display only the count of matching lines.
 - The **-n** option means show line number(s) and matching line(s)
 - The **-v** option means show all lines that do not match the search string.
 - For more options, look up the manual page (man grep)
- **sort:** sort the contents of a file alphabetically
 - The **-r** option means reverse sorting.
 - The **-n** option means sort numerically
 - The **-f** option means case insensitive sort.
 - The **-kN** option means sort by column N.

Piping Commands

- More examples:
- `du -h | sort -h | tail -10`
 - **du -h**: display disk usage for all folders recursively in the current folder *in human readable format (eg 1K, 3M, 5G)*
 - **sort -h**: sort lines of text in ascending order *comparing human readable numbers (eg 1K, 3M, 5G)*
 - **tail -10**: display the last part of a file or stream, *specifically the last 10 lines.*
- Thus: this command displays the top 10 folders consuming the most disk space.

Piping Commands

- More examples:
- `ls -al | grep ^d`
 - Run the “`ls -al`” command and pipe the output in to `grep`
 - **`grep ^d`**: match all lines beginning with the character ‘d’
 - Thus: the command displays detailed information about all subfolders (and only subfolders) in the current working directory

Piping Commands

- More examples:
- `ls -al | grep ^d -v | cut -c58-`
 - Run the “**ls -al**” command and pipe the output in to grep.
 - **grep ^d -v**: the **minus v** option inverts the sense of matching, thus: match every line that does not begin with a ‘d’ character.
 - **cut -58-**: the cut command prints everything from each line starting at column 58 until the end of the line.
- Thus: the command displays the names of the files (and only the files) in the current working directory.