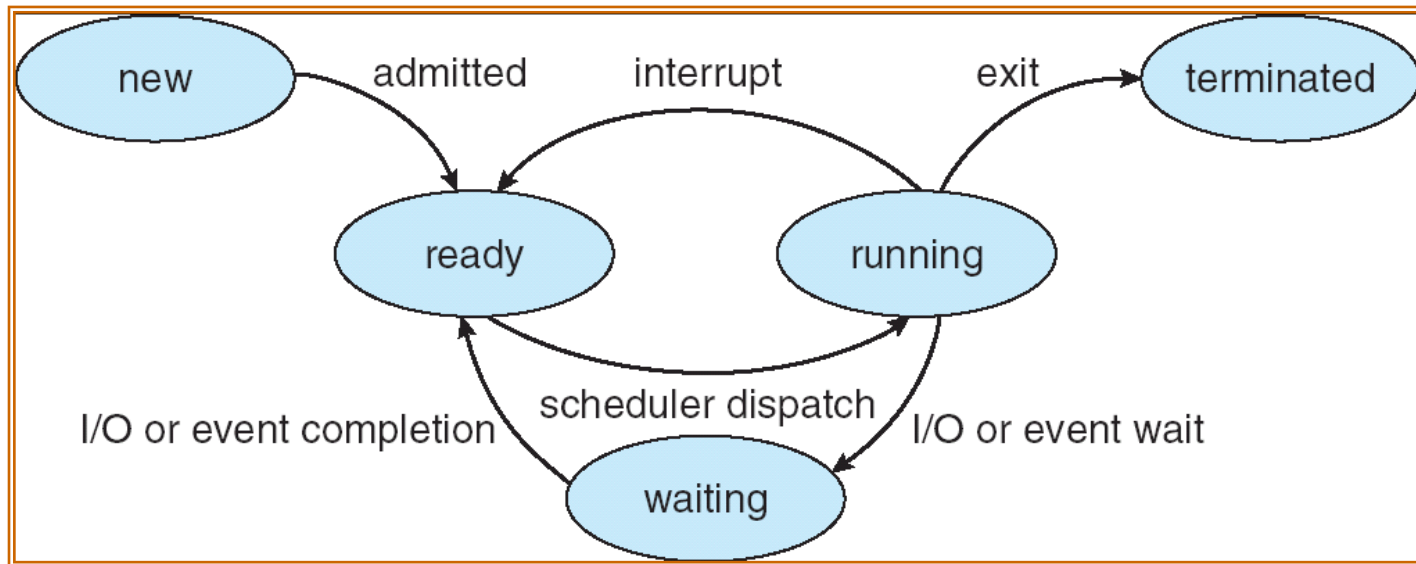


# Operating Systems Fundamentals

## Introduction to Processes



# Introduction

- A **Process** is an **application / program in execution**.

In other words, it is the execution of a sequence of instructions (program), whose net result is the provision of some system function or user function.

We say that:

*A processor runs a process,*

*or*

*A process runs on a processor.*

# Processes

- Processes are created whenever you start to run an application, i.e. when the operating System loads a program (e.g. a browser)

*A program becomes a process **when it is loaded into memory.***

- In the next few lectures we will look at the operations and aspects of the Operating System that support Processes and user applications

# Processes

---

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems


# Process Concept

- An OS executes a variety of programs & services
  - Batch system Jobs
    - e.g. cron jobs at night time, processing of payrolls, electricity bills  
not concerned with interacting with users , jobs are Q'ed up and run
  - Time-shared systems – User Programs, Tasks
    - Many exe's/windows running in parallel at same time
    - How does an OS do that??
- We know that a process is an **application / program in execution**
  - Process execution must progress in **sequential fashion**
    - Recall Assembly language and the Program Counter

# Process Concept, continued

- A Process must be self contained
    - All information a process needs to operate correctly ... needs to be under the process's control
  - To achieve this, a process includes:
    - Program Counter      Specifies the next instruction to execute
    - Stack                      Used to store temporary data
    - Text/Data Sections    This is the program code and global data
- Recall your computer architecture from the current semester

# A Process in Memory is more than the program code



The **process stack** contains temporary data, for example function parameters, return addresses and local variables.

The **heap** is memory that is dynamically allocated, i.e. allocated during run time

The **data section** contains global variables.

The **text** is the program code

# Question?

---

If a user invokes multiple copies of the same web browser,  
do you think this is the same process?



# Question?

If a user invokes multiple copies of the same web browser,  
do you think this is the same process?

## ANSWER

No - each is a **separate** process.

The **text sections** are the same

but the **data, heap and stack sections** are different for each process

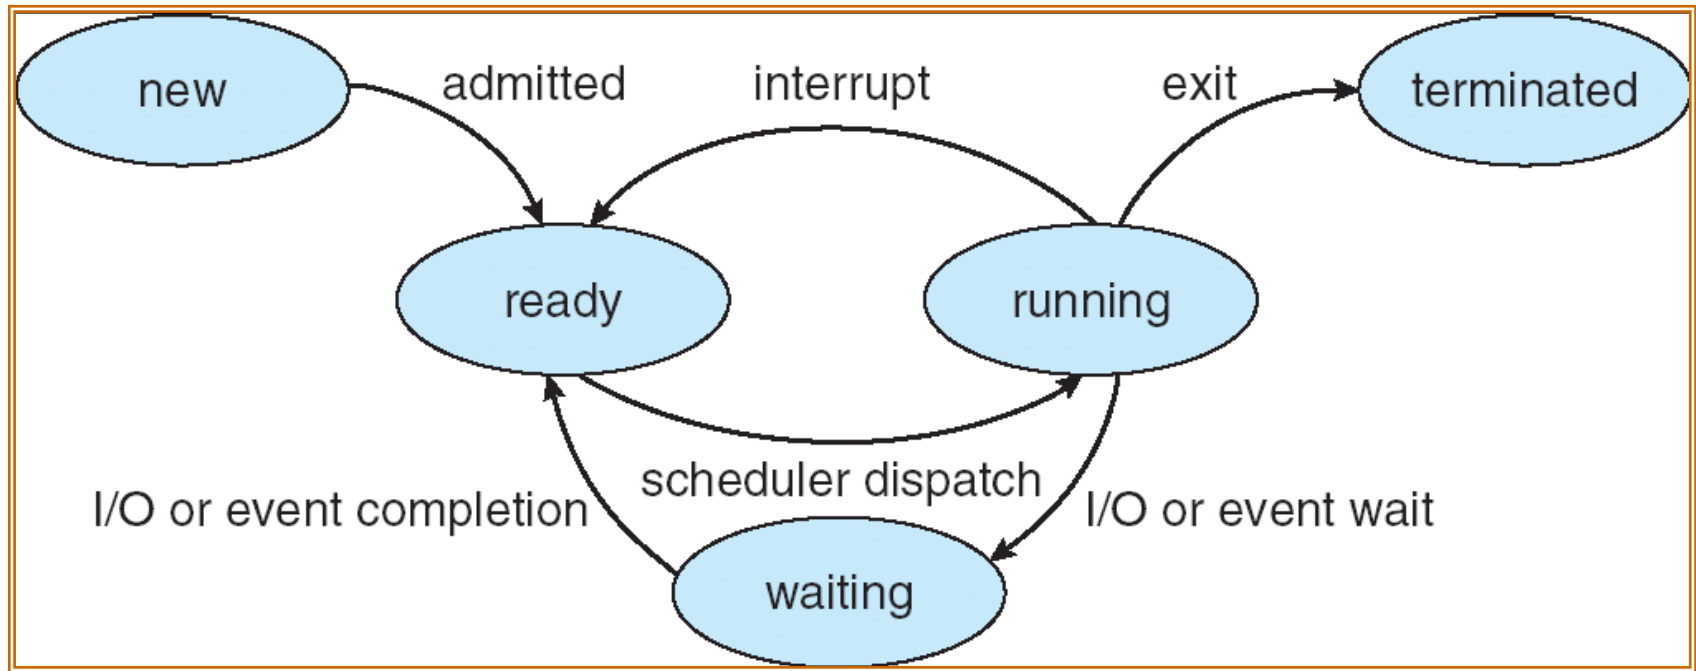
# Process States

- A Multi-tasking Operating System supports many processes
- Each process must be manageable and controllable
- To achieve this, processes have different **states** that control how they operate and behave

# Process State

- As a process executes, it changes *state*
  - **New:** The Process is being created / started.
  - **Running:** Instructions are being executed / run on the CPU.
  - **Waiting:** The Process is waiting for some event to occur. e.g. Hard disk access,  
Keyboard operation
  - **Ready:** The Process is waiting to be assigned to a CPU / It is ready to be run.
  - **Terminated:** The Process has finished execution.

# Process State Transitions



# Process State Transitions, continued

- Progression between the states is controlled by the Operating System
  - An **Interrupt** can be generated by two things
    - An external application/hardware event or
    - the time slice assigned by the scheduler expires
  - **I/O or wait** event arises if the process must pause to gather extra data.
    - As this can take some time, CPU is freed to allow other processes to run
  - **Scheduler Dispatch** means the process can be run on the CPU based on the OS schedule scheme  
(more on OS schedules later...)

# Process Data

- To allow an Operating System support many processes at once, each process needs specific data
- This data allows the OS to easily find out about each process without having to remember every process that is running
- The data is stored in a **Process Control Block**, also called **Process Descriptor**

# Process Control Block (PCB) Features

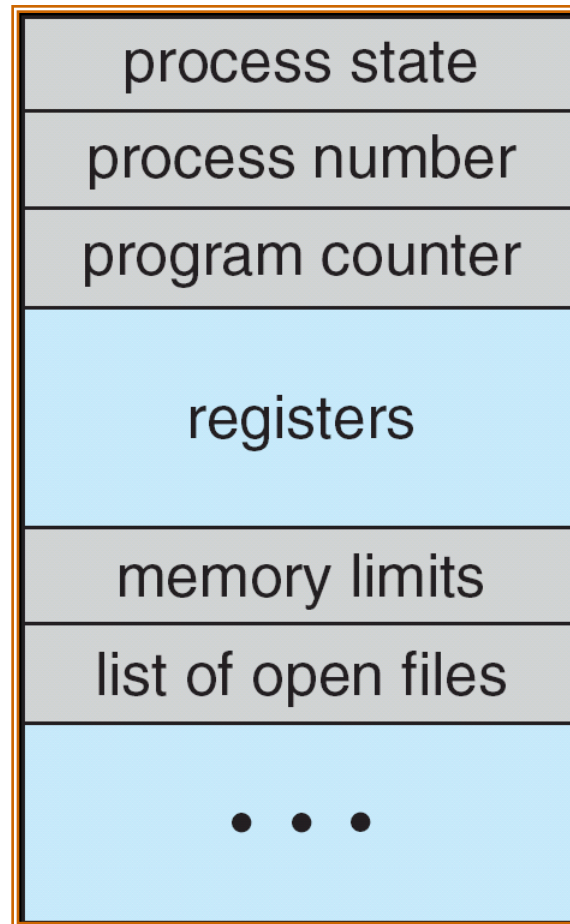
- The **Process Control Block** is an area in memory that stores the following process information that **uniquely characterises a particular process**.
  - Process state
  - Program counter
  - CPU register contents
  - CPU scheduling information
  - Memory-management information
  - Accounting information
  - I/O status information

# Process Control Block (PCB) Features

- Process state
  - *Ready, running, waiting, etc*
- Program counter
  - *Next instruction to be executed for that process*
- CPU register contents
  - *Values*
- CPU scheduling information
  - *Used by short-term scheduler/ CPU scheduler. Includes a process priority etc*
- Memory-management information
  - *Memory limits.*
- Accounting information
  - *e.g. How much CPU used, when it was last run, etc*
- I/O status information
  - *e.g. List of I/O devices allocated to process, list of open files...*



# PCB, continued



# Context Switch

A **context switch** is the switching of the CPU from one process (or thread) to another.

Also known as **CPU Switch**

# Context Switch

The OS kernel:

- stops *execution of the process* currently using the CPU.
- resumes *execution of some other process* that had previously been queued up waiting to use the CPU.

Process information stored in PCBs

# Context Switch

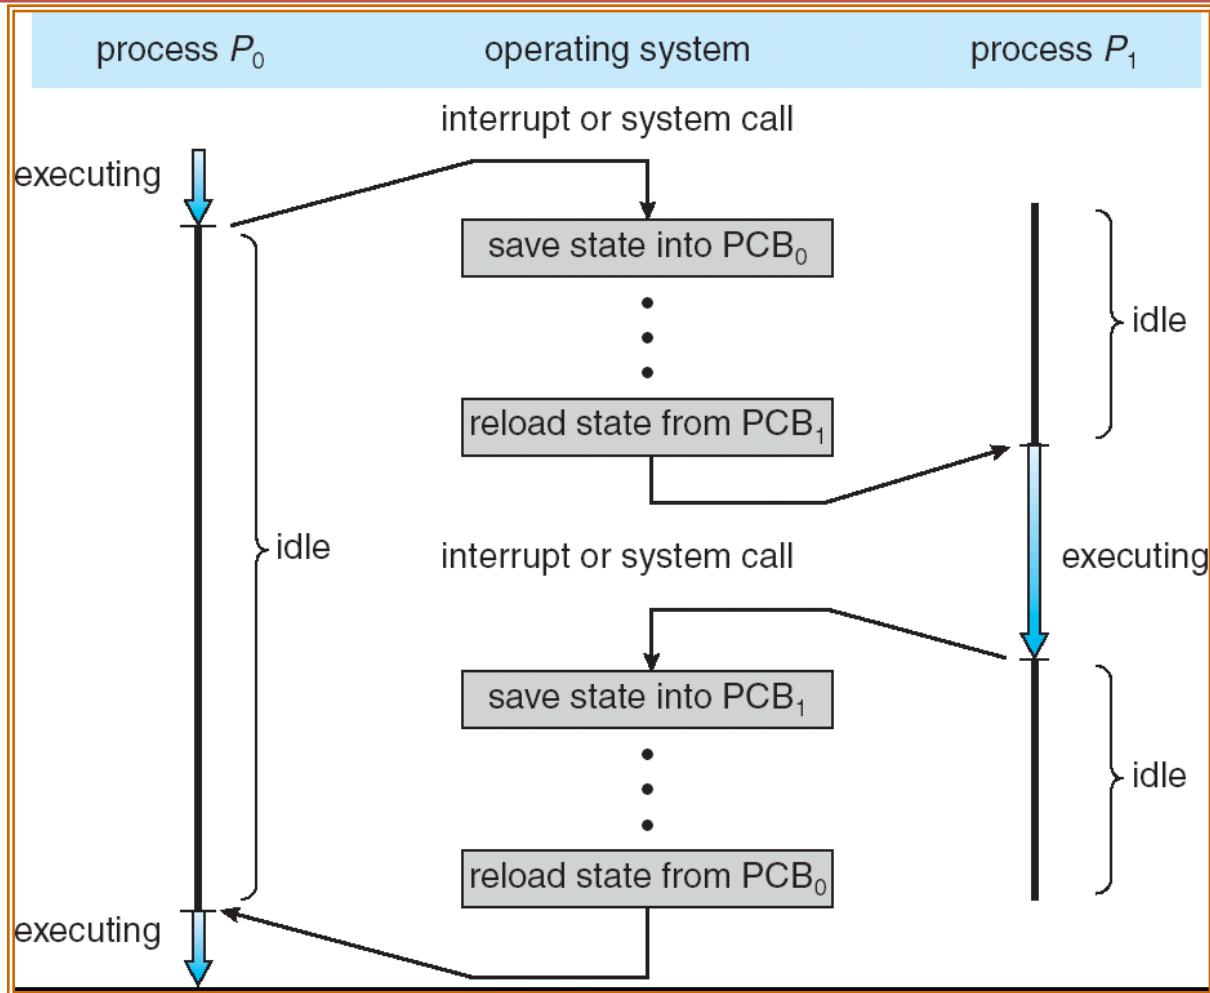
- It is not possible to read all registers and write all registers without some time overhead
  - The registers can be read quickly but only 1 at a time
  - There will be many registers to access and check
  - So .... This loses us time ... but ....

# Context Switch, continued

- The advantages of Multi-tasking **outweigh** the overheads of a context switch
  - Processors are getting faster while context switch operations are remaining basically the same
- The Process Control Block (PCB) supports the context switch by storing all its own process information
  - This means the OS does not need to store this information; it only needs to read and access the PCB information for a new process

# Context Switch, continued

## $P_0$ running, then $P_1$ , then $P_0$



# Process Scheduling

- The objective of **multiprogramming**, is to have **some process running at all times**. in order to maximise CPU utilisation
- The objective of **timesharing** is to switch the CPU among processes so frequently that users can interact with each program while running, without seeing any program stoppage or delay
- To meet these 2 objectives, the **process scheduler** selects an **available process** for program **execution** on the CPU.

# Process Queues

- As we know, at any one time, many different processes exist.
- For example, as we have seen, processes may be the following states
  - **Running** using the CPU
  - **Waiting** for I/O (input/output)
  - **Ready** to be run
- Depending on its state, a process is put into a particular type of **process queue**.




# Process Queues

- **Job queue:** Set of all processes in the system
  - e.g. Look at Task Manager,  
all processes running on the computer are shown
- **Ready queue:** Set of all processes residing in main memory and ready to be executed
- **Device queue:** Set of processes waiting for an I/O device
  - There are many device queues  
as each device will have its own queue
- As a process executes,  
it **migrates** among the various queues

# Process Queues

- As a process enters the system they are put into a **job** queue.
- The queues are generally stored as **linked lists**.
- Example of process migration:
  - A new process is put in the **READY** QUEUE.
  - Once the process is allocated the CPU and is executing, it could issue an I/O request and then be put into the **DEVICE( I/O)** QUEUE.

# Process Queues, continued



Example:

Linked list  
design

# Process Queues, continued

- The number of processes in each queue is important to the Operating System.
- It needs to know how much resources to apply to each
  - There is no point in giving time to operating the Device Queues if nothing is waiting for that device
- The choice of what to do with a Queue is handled by a **Scheduler**
  - We will be looking at schedulers later

# Summary

---

- The concept of a Process has been introduced
- The data a Process must hold and store in PCB has been shown
- Process States and the operations at each have been presented
- Notion of how processes move between various queues also detailed