

# Operating Systems Fundamentals

## Process Schedulers



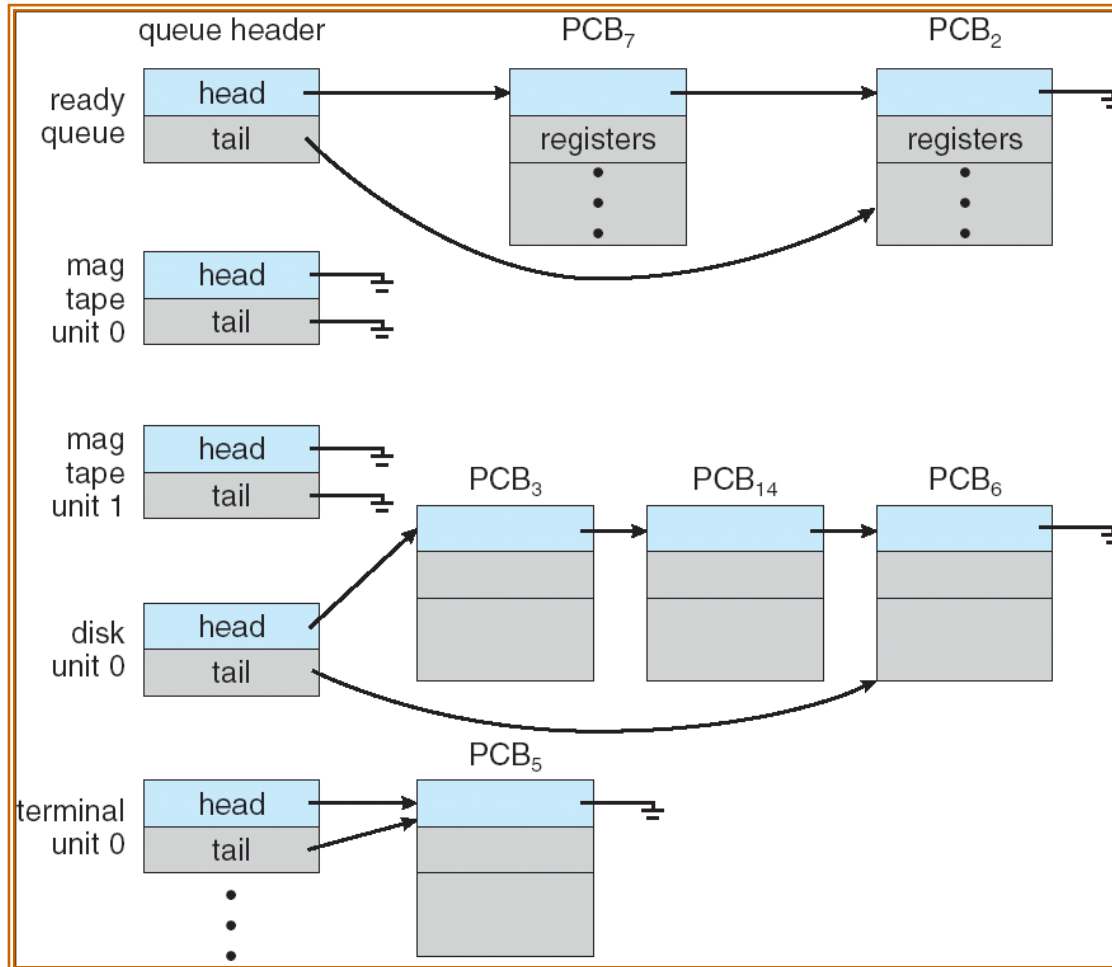
# Introduction

- Previously seen the idea of a **Process**  
“An application / program in Execution”
- Discussed the concepts of **process queues** and different **process states**
- **Next** we look briefly at how the Operating System **manages** the various **states** and **controls the execution** of many different processes

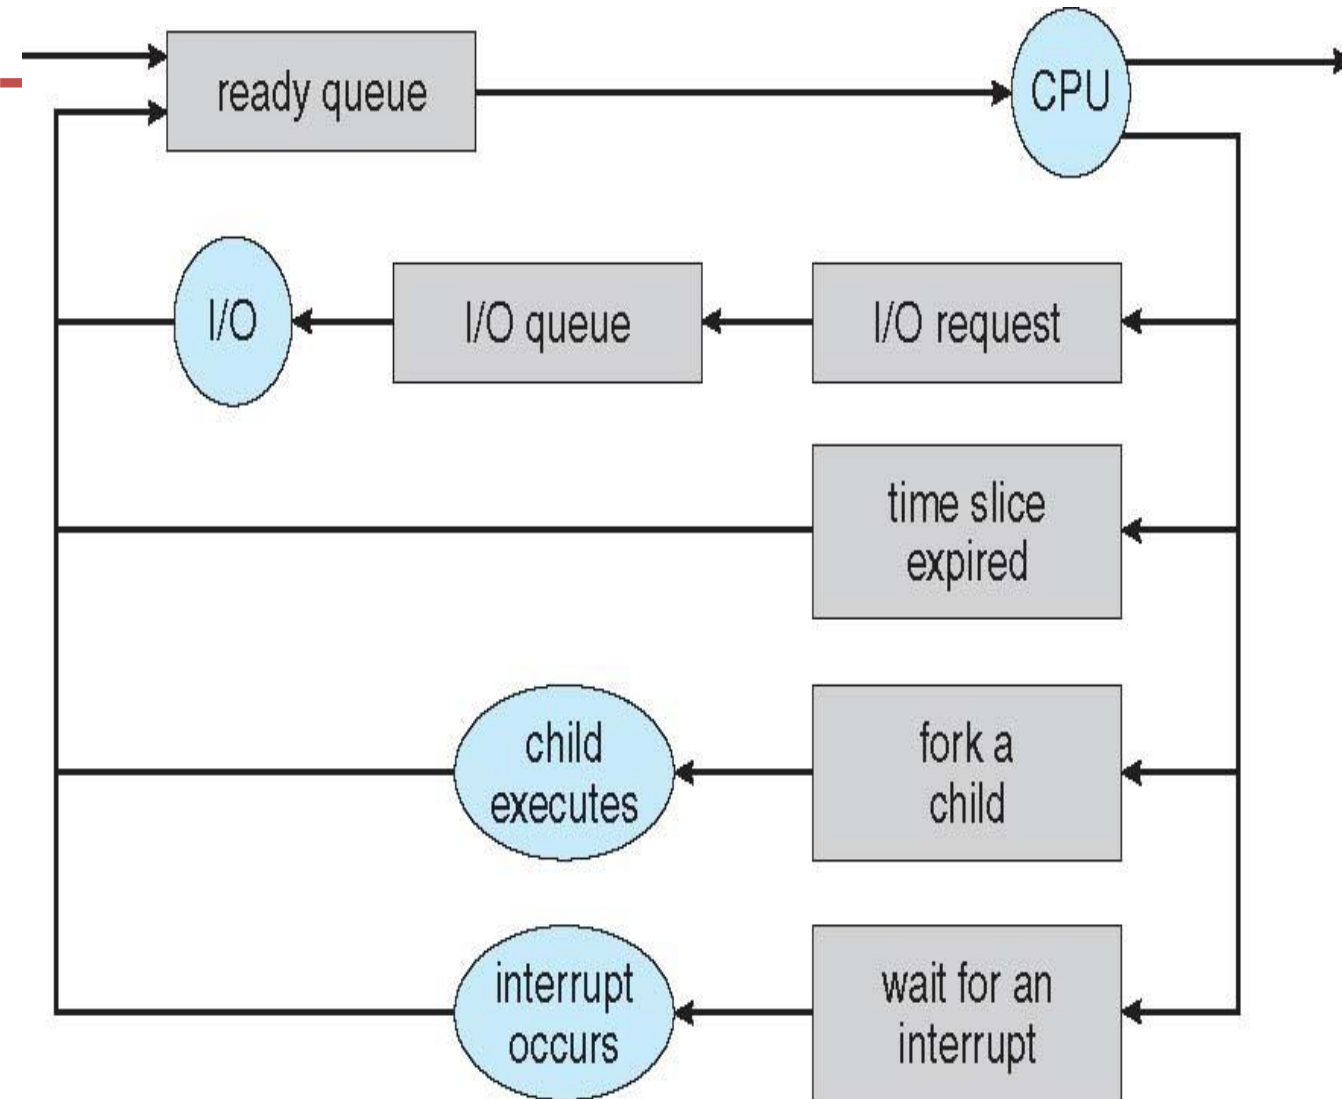
# Queues Recapped

- **Job Queue; Ready Queue; Device Queue**
- Processes
  - **migrate** between each queue
  - *depending on the state they are in*
- The Operating System is responsible for
  - **moving** processes between the queues
- There will be a **unique** queue
  - for each **hardware Device** present in the computer

# Recall, Queue Layout



# Process Migrating among queues



Processes move between the various queues depending on the operations they have to perform

When a process **terminates** – it is **removed** from all queues and has its PCB and resources de-allocated.

# Process Scheduling

- As well as the queues,  
  
the OS need a means to schedule
  - which process is to run and
  - which processes should move between queues
- Some processes are more important than others and need to run on the CPU more frequently
  - Have higher priority

# Degree of Multiprogramming

The “**degree of multiprogramming**” refers to the number of processes in memory.

If

the average rate of process creation

=

average departure rate of processes  
leaving the system

Then the ***degree of multiprogramming*** is **stable**

# Main Schedulers

- Two schedulers support the overall operations of the Processes

**Long-term** scheduler (or *Job* scheduler)

and

**Short-term** scheduler (or *CPU* scheduler)


The primary distinction between the schedulers lies in the **frequency of execution**



# Schedulers

**Long Term** feeds the Ready Q

**Short term** feeds the CPU



# Schedulers

- **Long-Term:**
  - Selects which processes should be brought into the ready queue
    - from new,
    - and from I/O (Device) Q.
  - Less frequently invoked than short term

# Schedulers

- **Short-Term**

- Selects which process should be executed on CPU next
- Moves processes between the Ready Queue and the CPU
- Moves processes from the CPU to the other appropriate queue after
  - its time slice is complete or
  - I/O event
- Short CPU time slices
- Very frequently invoked (milliseconds)

# Question?

**“The Long-term scheduler controls the degree of multiprogramming”**

What do we mean by this statement?

Recall: If average rate of process creation =  
average departure rate of processes leaving the system  
Then the *degree of multiprogramming* is stable.

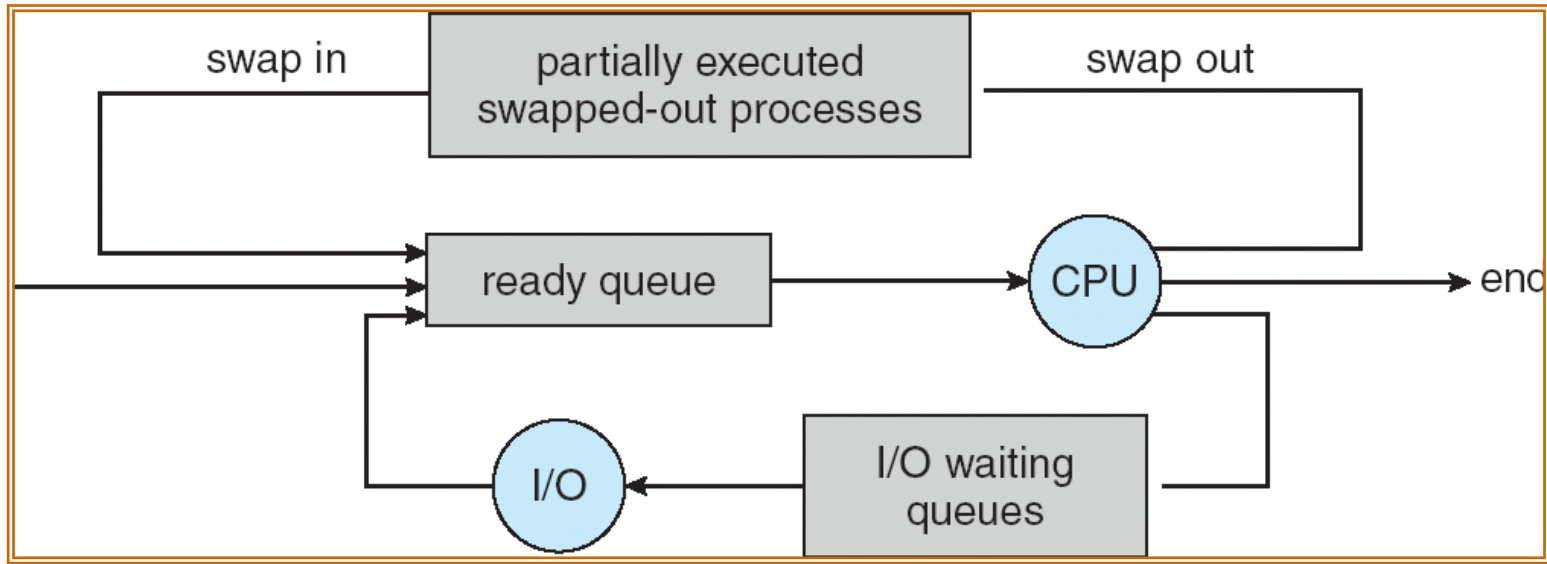
# Medium Term Scheduling

- Some Operating systems also have an additional, intermediate form of scheduling called the **Medium Term** Scheduler
  - The difference between the Long term scheduler and the Short term scheduler operations means we may not get the best overall performance or efficiency out of all processes.

# Medium Term Scheduling

- The medium term scheduler
  - removes **partially** executed process from fast memory
  - i.e. **swaps it out** to slower disk memory
- Later it brings the process back into memory
  - i.e. **swaps it back in** and executes from where it left off.
- It might “**swap out**” a process that:
  - has been inactive for a while
  - has a low-priority
  - is taking up a large amount of memory

# Medium Term Scheduling



## Advantages of Medium-Term Scheduler

- Reduces the degree of multiprogramming (reduces active contention for the CPU)
- Improves the process mix (i.e. I/o bound and CPU bound)
- Frees up memory when required.

SO, BOTH THE LONG-TERM and the MEDIUM-TERM SCHEDULER CONTROLS THE DEGREE OF MULTIPROGRAMMING

# Multitasking in Mobile Systems

- Some systems / early systems allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
  - Single **foreground** process- controlled via user interface
  - Multiple **background** processes– in memory, running, but not on the display, and with limits
  - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
  - Background process uses a **service** to perform tasks
  - Service can keep running even if background process is suspended
  - Service has no user interface, small memory use



# Process Category

---

**influenced by the way  
they interact with the schedulers**

# Process Category Summary

---

- **Computation** based processes (aka **CPU-Bound**)
- **I/O** based processes (aka **I/O-Bound**)
- **Balanced** processes
  - where neither of above are predominant in the operations

# Process Category Summary

- **Computation** based processes
  - Spends more time doing calculations.
  - Use all time slice for calculations
  - Uses fewer, but longer CPU bursts.
  - Heavy demand on CPU.
  - All data will typically be present in
    - registers or
    - caches and
    - can be accessed quickly

# Process Category Summary

- **I/O** based processes
  - A lot of operations are based on accessing
    - user data
    - hardware resources.
  - Process usually runs a few instructions before
    - it must access an external device or
    - wait for answers from the user
  - Therefore ... many short CPU bursts.
- Recall speed difference between CPU and external I/O devices
  - 1 second waiting for user responses would allow a large number of processes to run

# Scheduler concerns

---

**It is important that the long-term scheduler selects a good mix of I/O-bound and CPU-bound processes**

**WHY?**

# Scheduler concerns

The long-term scheduler must select a good mix of I/O bound and CPU-bound processes because:

- If all processes are I/O bound then the ready queue will almost always be empty and the short-term scheduler will have little to do.
- 1) If all processes are CPU bound, the I/O waiting queue will be empty, devices will go unused, and the system will be unbalanced

The system with the best performance will have the best mix of CPU-bound and I/O bound processes.

# Operations on Processes – Process Creation and Termination

- The processes in most systems can execute **concurrently**
- These systems must provide a mechanism for
  - Process **creation** and
  - Process **termination**.

# Why/When are Processes Created?

## Examples:

- **Interactive logon**

A user at a terminal logs on to the system.

- **User starts a program**

- **Created by OS to provide a service**

The OS can create a process to perform a function on behalf of a user program, without the user having to wait (e.g., a process to control printing).

- **Spawned by existing process**

For purposes of modularity or to exploit parallelism, a user program can dictate the creation of a number of processes.  
e.g. chrome tabs



# Process Creation

- A process may create several new processes,
  - via a create-process system call,
  - during the course of execution.
- The creating process is called a **parent** process,
  - and new processes are called the **children** of that process
- Parent process create children processes,
  - which, in turn create other processes, forming a **tree** of processes.
- **Two execution** options
  - Parent and children execute **concurrently**
  - Parent **waits** until children terminate

# Process Resource Sharing

- Processes need resources to accomplish their task (CPU time, files, memory, I/O devices).
- There are **three** ways that parent and child processes share resources:
  1. Parent and children processes **share all** resources
  2. Children share a subset of parent's resources
    - this helps prevent overload of system by creating too many sub-processes
  3. Parent and children share no resources

# Process Termination

- Processes eventually
  - complete
  - or are terminated
- There are two approaches to termination
  - The Process itself executes last statement
    - and asks the operating system to delete it (**exit**)
  - Parent
    - may terminate execution of children processes (**abort**)

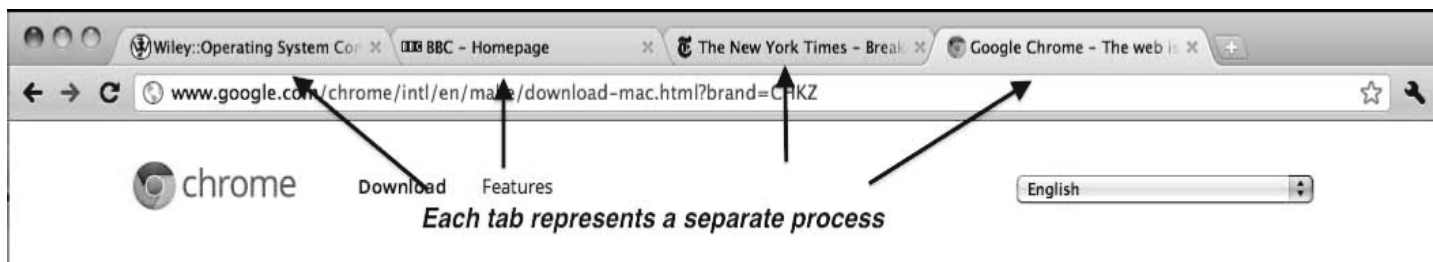
# Possible Reasons for process termination

---

- Normal completion
- Arithmetic error, or data misuse (e.g., wrong type)
- Invalid instruction execution
- Insufficient memory available, or memory bounds violation
- Resource protection error
- I/O failure

# Multiprocess Architecture – Chrome Browser

- Many web browsers ran as single process (some still do)
  - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multi-process with 3 categories
  - **Browser** process manages user interface, disk and network I/O
  - **Renderer** process
    - Renders web pages, deals with HTML, Javascript,
    - New one for each website opened
    - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
  - **Plug-in** process for each type of plug-in



# Chrome processes

- To examine the processes inside of Chrome
  - Menu-> MoreTools -> Task Manager

Task Manager - Google Chrome

Task	Memory	CPU	Network	Process ID
• Browser	101,964K	3.9	0	3760
• GPU Process	49,924K	1.5	0	1936
• Tab: Course: Operating Systems Fundamentals	31,964K	0.0	0	5396
• <b>Tab: Google</b>	<b>24,236K</b>	<b>0.0</b>	<b>0</b>	<b>2532</b>
• Extension: Cisco WebEx Extension	1,632K	0.0	0	4592
• Extension: Office Editing for Docs, Sheets & Slides	46,428K	0.0	0	5876
• Background Page: Videostream for Google Chromecast™	24,264K	0.0	0	4672

End process

# Example of UNIX Process creation

---

- Two styles of processes can be created in Unix
  - Parent Process
  - Child Process

# Unix process Creation – child and parent


- When a process is started, a duplicate of that process is created.
  - This new process is called the **child** and the process that created it is called the **parent**.
  - The child process then replaces the copy for the code the parent process created with the code the child process must execute.



# UNIX Process Creation

1. Fork() system call in UNIX creates a new (child) process

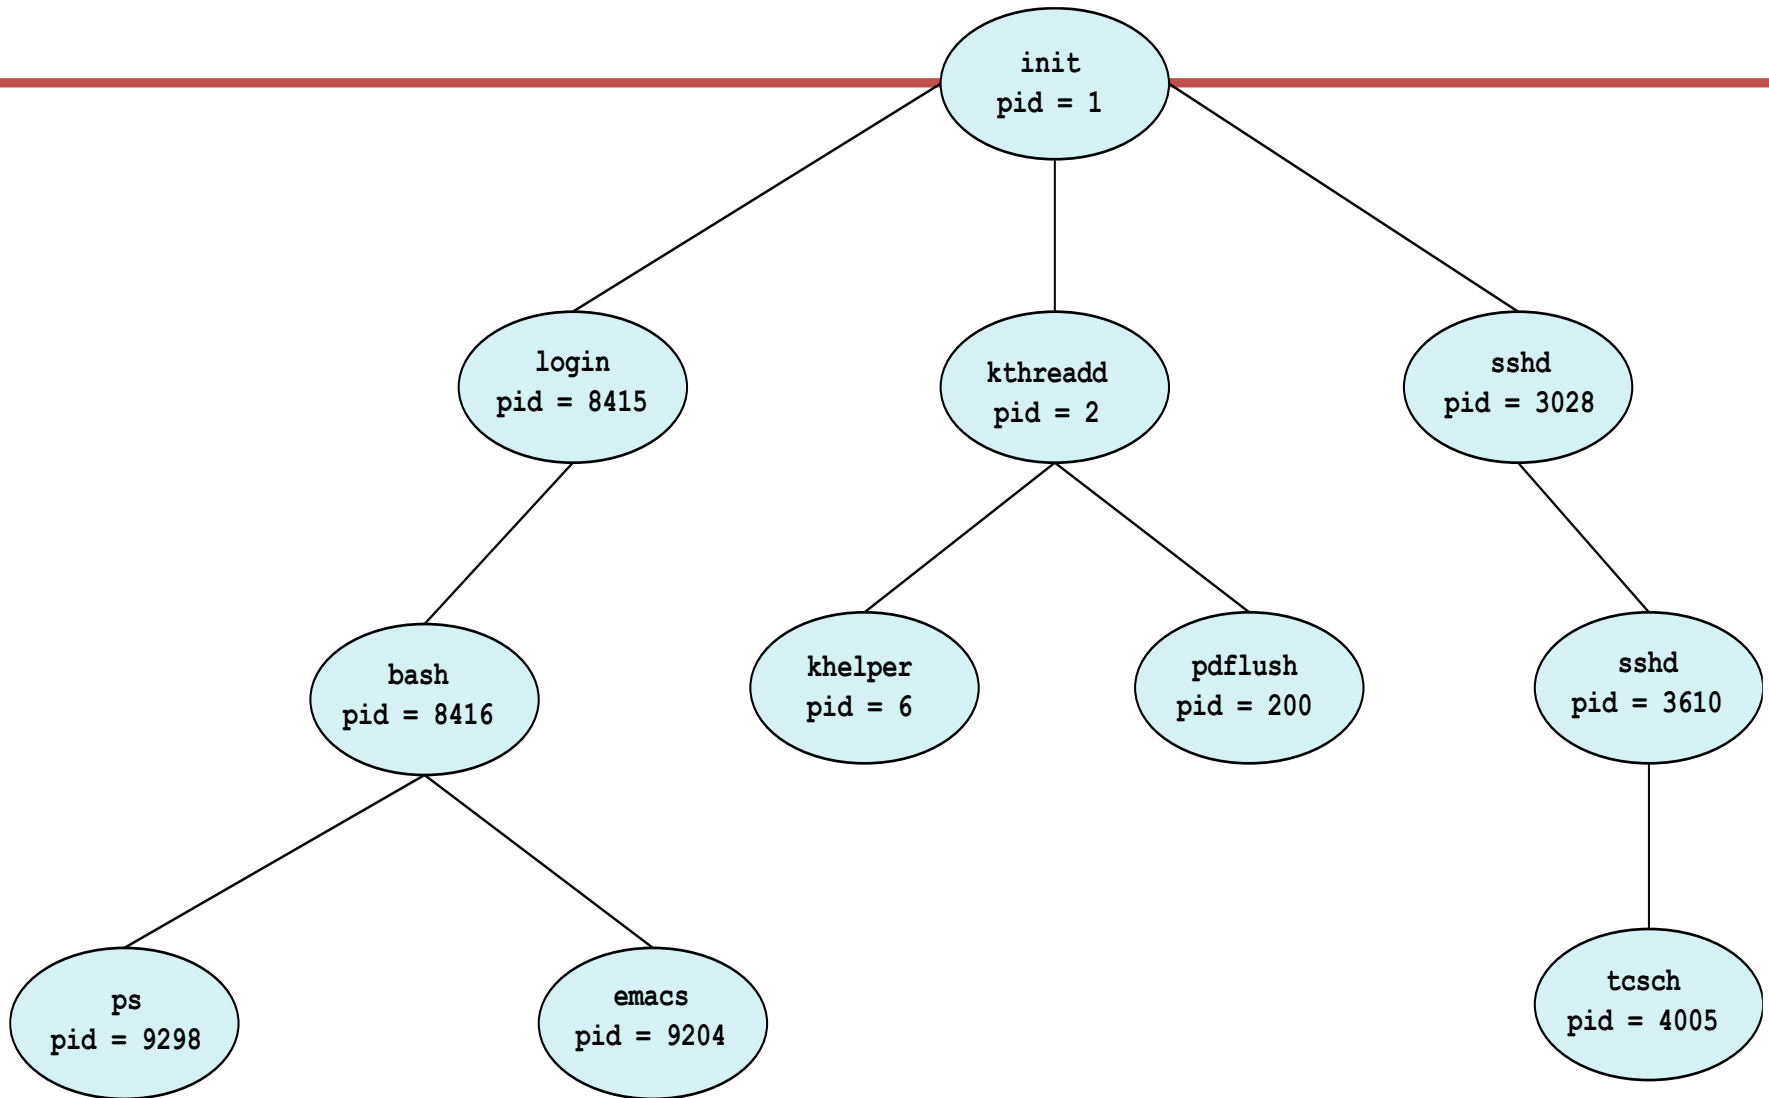
2. Exec() system call used after a fork() system call to replace the memory space with a new program



# UNIX Parent, Child Relation

- All processes are created as children of the Operating System
- The children are then converted to Parent processes which can create other Child processes
- This repeats until all processes and services of the OS are running
- This creates a tree / hierarchy of processes
  - Also creates families as similar processes take information from their parent

# A tree of processes in Linux



# UNIX Process Termination

- After completing work a process may execute **exit()** system call
  - asking the OS to delete it
    - parent is notified
    - process' resources are de-allocated by operating system
- Parent may terminate execution of a child process - **abort()** system call possibly because:
  - task assigned to child is no longer required
  - child exceeded allocated resources
- if parent exits (some) OS's require all children to terminate – cascading termination

# Summary

- Different Process Schedulers have been introduced
- Different styles of processes have been discussed
  - CPU-Bound, I/O Bound
- Idea of Process Creation and Process Termination also discussed
  - Idea of Parent and Child process